

FEDERAL STATE AUTONOMOUS EDUCATIONAL INSTITUTION
OF HIGHER EDUCATION
ITMO UNIVERSITY

Parallel algorithms for the analysis and synthesis of data
on the assignments No.9, 10, 11


Performed by
Ivan Dubinin
J4132c

St. Petersburg
2021

Assignment 9

Write an MPI program in which the global vector addition operation is modeled by adoubling (cascade) scheme using point-to-point data transfers. Compare the execution time of such a simulation using the MPI_Reduce procedure on as many processes as possible. Each process stores an array of 1,000,000 elements equal to '1'. The resulting array is saved, for example, on process 0. To check the result, print the first 5 elements for both methods to the console.

For implementing cascade scheme the following distribution of data was used. In my case a_i stands for array parts (+ is vector sum in this case)



Listing of the program

[Cascade variant](#)

[Reduce variant](#)

Description

The first program implements finding array SUM using MPI_Reduce call. The second one implements cascade method, shown on the picture above. Both programs were tested four times with eight procs. The results of this testing is shown on the pictures below. As one can see, time calculations for both approaches are not stable, but all in all for Reduce variant the average execution is 30 percent less than for Cascade one.

Example of launch parameters and output

```

[pes@vandosik HW_MPI]$ mpic++ Assignment9_reduce.cpp -o reduce_9
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus reduce_9 --mca opal_warn_on_missing_libcudat 0
Root process #0: final answer is : 1000000
Point to point answer merging took 134 us
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus reduce_9 --mca opal_warn_on_missing_libcudat 0
Root process #0: final answer is : 1000000
Point to point answer merging took 129 us
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus reduce_9 --mca opal_warn_on_missing_libcudat 0
Root process #0: final answer is : 1000000
Point to point answer merging took 14 us
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus reduce_9 --mca opal_warn_on_missing_libcudat 0
Root process #0: final answer is : 1000000
Point to point answer merging took 16 us

```

```

[pes@vandosik HW_MPI]$ mpic++ Assignment9_cascade.cpp -o cascade_9
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus cascade_9 --mca opal_warn_on_missing_libcudat 0
Reduce order 2, Process #4: received sum 125000 from #5, total sum is 250000
Reduce order 4, Process #4: received sum 250000 from #6, total sum is 500000
Reduce order 2, Process #6: received sum 125000 from #7, total sum is 250000
Reduce order 2, Process #0: received sum 125000 from #1, total sum is 250000
Reduce order 4, Process #0: received sum 500000 from #4, total sum is 750000
Reduce order 8, Process #0: received sum 250000 from #2, total sum is 1000000
Root process #0: final answer is : 1000000
Point to point answer merging took 23 us
Reduce order 2, Process #2: received sum 125000 from #3, total sum is 250000
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus cascade_9 --mca opal_warn_on_missing_libcudat 0
Reduce order 2, Process #0: received sum 125000 from #1, total sum is 250000
Reduce order 2, Process #2: received sum 125000 from #3, total sum is 250000
Reduce order 2, Process #4: received sum 125000 from #5, total sum is 250000
Reduce order 4, Process #4: received sum 250000 from #6, total sum is 500000
Reduce order 2, Process #6: received sum 125000 from #7, total sum is 250000
Reduce order 4, Process #0: received sum 250000 from #2, total sum is 500000
Reduce order 8, Process #0: received sum 500000 from #4, total sum is 1000000
Root process #0: final answer is : 1000000
Point to point answer merging took 70 us
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus cascade_9 --mca opal_warn_on_missing_libcudat 0
Reduce order 2, Process #0: received sum 125000 from #1, total sum is 250000
Reduce order 2, Process #2: received sum 125000 from #3, total sum is 250000
Reduce order 2, Process #4: received sum 125000 from #5, total sum is 250000
Reduce order 4, Process #0: received sum 250000 from #2, total sum is 500000
Reduce order 2, Process #6: received sum 125000 from #7, total sum is 250000
Reduce order 4, Process #4: received sum 250000 from #6, total sum is 500000
Reduce order 8, Process #0: received sum 500000 from #4, total sum is 1000000
Root process #0: final answer is : 1000000
Point to point answer merging took 275 us
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus cascade_9 --mca opal_warn_on_missing_libcudat 0
Reduce order 2, Process #6: received sum 125000 from #7, total sum is 250000
Reduce order 2, Process #0: received sum 125000 from #1, total sum is 250000
Reduce order 4, Process #0: received sum 250000 from #2, total sum is 500000
Reduce order 8, Process #0: received sum 500000 from #4, total sum is 1000000
Root process #0: final answer is : 1000000
Point to point answer merging took 37 us
Reduce order 2, Process #2: received sum 125000 from #3, total sum is 250000
Reduce order 2, Process #4: received sum 125000 from #5, total sum is 250000
Reduce order 4, Process #4: received sum 250000 from #6, total sum is 500000
[pes@vandosik HW_MPI]$

```

Assignment 10

Complete the program Assignment10.c. Compile and run it.

Study the code carefully and explain how it works.

Listing of the program

[See it in my github repo](#)

Description

Each proc sends its rank to previous and next procs. Non-blocking MPI_Irecv and MPI_Isend calls are used.

All requests completion is determined by calling the MPI_Waitall.

Example of launch parameters and output

```
[pes@vandosik HW_MPI]$ mpic++ Assignment10.c -o task_10
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus task_10 --mca opal_warn_on_missing_libcudatoolkit 0
Process: 0 Prev_data: 7 Next data: 1
Process: 1 Prev_data: 0 Next data: 2
Process: 2 Prev_data: 1 Next data: 3
Process: 3 Prev_data: 2 Next data: 4
Process: 4 Prev_data: 3 Next data: 5
Process: 5 Prev_data: 4 Next data: 6
Process: 6 Prev_data: 5 Next data: 7
Process: 7 Prev_data: 6 Next data: 0
[pes@vandosik HW_MPI]$
```

Assignment 8

Based on Assignment 10, write a program for ring topology exchange using the `MPI_Sendrecv()` function. In situations where you need to exchange data between processes, it is safer to use the overlaid `MPI_Sendrecv` operation. The `MPI_Sendrecv` function combines the execution of the send and receive operations. Both operations use the same communicator, but message IDs may differ. The location of the received and transmitted data in the address space of the process should not overlap. The data sent can be of different types and lengths. In cases when it is necessary to exchange data of the same type with replacement of the sent data with the received ones, it is more convenient to use the `MPI_Sendrecv_replace` function. In this operation, the data sent from the buf array is replaced with the received data. The special address `MPI_PROC_NULL` can be used for source and dest in data transfer operations. Communication operations with such an address do nothing. The use of this address is convenient instead of using logical constructs to analyze the conditions to send / read a message or not.

Listing of the program

[See it in my github repo](#)

Description

Program created n processes. Each process receives and sends message, equals its rank to previous and next (relatively its rank) neighbor. Main difference in new program – usage of `MPI_Sendrecv` function to handle sending and receiving messages. In first `MPI_Sendrecv` operation every process sends message to previous and receives from next. In second `MPI_Sendrecv` operation every process sends message to next and receives from previous.

Example of launch parameters and output

```
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus task_11 --mca opal_warn_on_missing_libcuda 0
Proccess: 3 Prev_data: 2 Next data: 4
Proccess: 1 Prev_data: 0 Next data: 2
Proccess: 2 Prev_data: 1 Next data: 3
Proccess: 4 Prev_data: 3 Next data: 5
Proccess: 5 Prev_data: 4 Next data: 6
Proccess: 6 Prev_data: 5 Next data: 7
Proccess: 7 Prev_data: 6 Next data: 0
Proccess: 0 Prev_data: 7 Next data: 1
[pes@vandosik HW_MPI]$ mpirun -n 8 --use-hwthread-cpus task_11 --mca opal_warn_on_missing_libcuda 0
Proccess: 2 Prev_data: 1 Next data: 3
Proccess: 3 Prev_data: 2 Next data: 4
Proccess: 4 Prev_data: 3 Next data: 5
Proccess: 5 Prev_data: 4 Next data: 6
Proccess: 6 Prev_data: 5 Next data: 7
Proccess: 0 Prev_data: 7 Next data: 1
Proccess: 1 Prev_data: 0 Next data: 2
Proccess: 7 Prev_data: 6 Next data: 0
[pes@vandosik HW_MPI]$
```